

# Investigate\_a\_Dataset

June 7, 2022

## 1 Project: Investigate a Dataset - Tmdb-Movies Dataset

### 1.1 Table of Contents

Introduction

Data Wrangling

Exploratory Data Analysis

Conclusions

## Introduction

#### 1.1.1 Dataset Description

: This is a dataset containing 10,000 movies, collated from 1960 - 2015. Collected from The Movie Database, it contains 21 columns . They are, \* the ID column - containing unique identifiers for each individual movie entry, \* Imdb Id - containing identifiers for each movie's result on Imdb, \* popularity - of each movie given a numerical value, \* budget - the budget for each movie, \* revenue - the revenue for each movie upon release, \* original\_title - the title of each individual movie, \* cast - the main cast of each movie, \* homepage - the link to each movie's release website, \* director - the director of each movie, \* tagline - the tagline or slogan given to each movie on release, \* keywords - the combination of keywords found in the movie, \* overview - an overview of the movie's description in string literal, \* runtime - the entire length of each movie in minutes, \* genres - the genre or genres under which each movie was released, \* production\_companies - the company behind each movie, \* release\_date - the date each movie was released, \* vote\_count, \* vote\_average, \* release\_year- the year in which each movie was released, \* budget\_adj - the budget of each movie adjusted for inflation in 2010 dollars, \* revenue\_adj - the revenue of each movie adjusted for inflation in 2010 dollars

#### 1.1.2 Question(s) for Analysis

Which Genres are the most popular from year to year? What kind of properties are associated with movies that have high revenue?

```
In [1]: # We import modules for use
import pandas as pd
import numpy as np
import ast
```

```
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

```
In [1]: # Upgrade pandas to use dataframe.explode() function.
!pip install --upgrade pandas==0.25.0
```

Collecting pandas==0.25.0

Downloading https://files.pythonhosted.org/packages/1d/9a/7eb9952f4b4d73fbd75ad1d5d6112f407e69  
100% || 10.5MB 3.2MB/s eta 0:00:01 4% | | 450kB 8.6MB/s eta 0:

Collecting numpy>=1.13.3 (from pandas==0.25.0)

Downloading https://files.pythonhosted.org/packages/45/b2/6c7545bb7a38754d63048c7696804a0d9473  
100% || 13.4MB 2.8MB/s eta 0:00:01 12% | | 1.6MB 24.3MB/s eta 0:

Requirement already satisfied, skipping upgrade: python-dateutil>=2.6.1 in /opt/conda/lib/python

Requirement already satisfied, skipping upgrade: pytz>=2017.2 in /opt/conda/lib/python3.6/site-p

Requirement already satisfied, skipping upgrade: six>=1.5 in /opt/conda/lib/python3.6/site-packa

tensorflow 1.3.0 requires tensorflow-tensorboard<0.2.0,>=0.1.0, which is not installed.

Installing collected packages: numpy, pandas

Found existing installation: numpy 1.12.1

Uninstalling numpy-1.12.1:

Successfully uninstalled numpy-1.12.1

Found existing installation: pandas 0.23.3

Uninstalling pandas-0.23.3:

Successfully uninstalled pandas-0.23.3

Successfully installed numpy-1.19.5 pandas-0.25.0

## ## Data Wrangling

### 1.1.3 General Properties

Here, we need to read in our data from a csv file containing the table for analysis. Our file is separated by a comma delimiter, so we do not need to specify a separator type. Afterwards, we inspect the data that has been read in by getting some info, some summary statistics and determining the shape and size of our data. This information will be useful in cleaning our data.

```
In [2]: # Load your data and print out a few lines. Perform operations to inspect data
# types and look for instances of missing or possibly errant data.
df = pd.read_csv('tmdb-movies.csv')
df.head()
```

```
Out[2]:
```

	id	imdb_id	popularity	budget	revenue	\
0	135397	tt0369610	32.985763	150000000	1513528810	
1	76341	tt1392190	28.419936	150000000	378436354	
2	262500	tt2908446	13.112507	110000000	295238201	
3	140607	tt2488496	11.173104	200000000	2068178225	
4	168259	tt2820852	9.335014	190000000	1506249360	

	original_title \		
0	Jurassic World		
1	Mad Max: Fury Road		
2	Insurgent		
3	Star Wars: The Force Awakens		
4	Furious 7		

	cast \		
0	Chris Pratt Bryce Dallas Howard Irrfan Khan Vi...		
1	Tom Hardy Charlize Theron Hugh Keays-Byrne Nic...		
2	Shailene Woodley Theo James Kate Winslet Ansel...		
3	Harrison Ford Mark Hamill Carrie Fisher Adam D...		
4	Vin Diesel Paul Walker Jason Statham Michelle ...		

	homepage	director \
0	<a href="http://www.jurassicworld.com/">http://www.jurassicworld.com/</a>	Colin Trevorrow
1	<a href="http://www.madmaxmovie.com/">http://www.madmaxmovie.com/</a>	George Miller
2	<a href="http://www.thedivergentseries.movie/#insurgent">http://www.thedivergentseries.movie/#insurgent</a>	Robert Schwentke
3	<a href="http://www.starwars.com/films/star-wars-episod...">http://www.starwars.com/films/star-wars-episod...</a>	J.J. Abrams
4	<a href="http://www.furious7.com/">http://www.furious7.com/</a>	James Wan

	tagline ... \
0	The park is open. ...
1	What a Lovely Day. ...
2	One Choice Can Destroy You ...
3	Every generation has a story. ...
4	Vengeance Hits Home ...

	overview runtime \
0	Twenty-two years after the events of Jurassic ... 124
1	An apocalyptic story set in the furthest reach... 120
2	Beatrice Prior must confront her inner demons ... 119
3	Thirty years after defeating the Galactic Empi... 136
4	Deckard Shaw seeks revenge against Dominic Tor... 137

	genres \
0	Action Adventure Science Fiction Thriller
1	Action Adventure Science Fiction Thriller
2	Adventure Science Fiction Thriller
3	Action Adventure Science Fiction Fantasy
4	Action Crime Thriller

	production_companies	release_date	vote_count \
0	Universal Studios Amblin Entertainment Legenda...	6/9/15	5562
1	Village Roadshow Pictures Kennedy Miller Produ...	5/13/15	6185
2	Summit Entertainment Mandeville Films Red Wago...	3/18/15	2480
3	Lucasfilm Truenorth Productions Bad Robot	12/15/15	5292
4	Universal Pictures Original Film Media Rights ...	4/1/15	2947

	vote_average	release_year	budget_adj	revenue_adj
0	6.5	2015	1.379999e+08	1.392446e+09
1	7.1	2015	1.379999e+08	3.481613e+08
2	6.3	2015	1.012000e+08	2.716190e+08
3	7.5	2015	1.839999e+08	1.902723e+09
4	7.3	2015	1.747999e+08	1.385749e+09

[5 rows x 21 columns]

```
In [3]: # Check through to determine how many empty values there are in the dataset
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10866 entries, 0 to 10865
Data columns (total 21 columns):
id                10866 non-null int64
imdb_id           10856 non-null object
popularity        10866 non-null float64
budget            10866 non-null int64
revenue           10866 non-null int64
original_title    10866 non-null object
cast              10790 non-null object
homepage          2936 non-null object
director          10822 non-null object
tagline           8042 non-null object
keywords          9373 non-null object
overview          10862 non-null object
runtime           10866 non-null int64
genres            10843 non-null object
production_companies 9836 non-null object
release_date      10866 non-null object
vote_count        10866 non-null int64
vote_average      10866 non-null float64
release_year      10866 non-null int64
budget_adj        10866 non-null float64
revenue_adj       10866 non-null float64
dtypes: float64(4), int64(6), object(11)
memory usage: 1.7+ MB
```

```
In [4]: # Determine number of rows and columns
df.shape
```

```
Out[4]: (10866, 21)
```

```
In [5]: # Summary statistics
df.describe()
```

```

Out[5]:
      id  popularity  budget  revenue  runtime \
count 10866.000000 10866.000000 1.086600e+04 1.086600e+04 10866.000000
mean  66064.177434   0.646441 1.462570e+07 3.982332e+07 102.070863
std   92130.136561   1.000185 3.091321e+07 1.170035e+08 31.381405
min     5.000000   0.000065 0.000000e+00 0.000000e+00 0.000000
25%   10596.250000   0.207583 0.000000e+00 0.000000e+00 90.000000
50%   20669.000000   0.383856 0.000000e+00 0.000000e+00 99.000000
75%   75610.000000   0.713817 1.500000e+07 2.400000e+07 111.000000
max   417859.000000 32.985763 4.250000e+08 2.781506e+09 900.000000

      vote_count  vote_average  release_year  budget_adj  revenue_adj
count 10866.000000 10866.000000 10866.000000 1.086600e+04 1.086600e+04
mean   217.389748   5.974922   2001.322658 1.755104e+07 5.136436e+07
std    575.619058   0.935142    12.812941 3.430616e+07 1.446325e+08
min     10.000000   1.500000   1960.000000 0.000000e+00 0.000000e+00
25%     17.000000   5.400000   1995.000000 0.000000e+00 0.000000e+00
50%     38.000000   6.000000   2006.000000 0.000000e+00 0.000000e+00
75%    145.750000   6.600000   2011.000000 2.085325e+07 3.369710e+07
max    9767.000000   9.200000   2015.000000 4.250000e+08 2.827124e+09

```

To begin our data cleaning process, we first take a look at how many empty rows there are in our data, then the number of duplicates, and finally the number of unique values.

```

In [6]: # Return sum of null values in dataset
df.isna().sum()

```

```

Out[6]: id                0
imdb_id                 10
popularity              0
budget                 0
revenue                0
original_title          0
cast                   76
homepage               7930
director               44
tagline               2824
keywords              1493
overview                4
runtime                0
genres                 23
production_companies   1030
release_date           0
vote_count             0
vote_average           0
release_year           0
budget_adj             0
revenue_adj           0
dtype: int64

```

```
In [7]: # Sum of duplicates
df.duplicated().sum()
```

```
Out[7]: 1
```

```
In [8]: df.nunique()
```

```
Out[8]: id                10865
imdb_id                10855
popularity            10814
budget                 557
revenue               4702
original_title        10571
cast                 10719
homepage              2896
director              5067
tagline               7997
keywords              8804
overview             10847
runtime               247
genres                2039
production_companies  7445
release_date          5909
vote_count            1289
vote_average           72
release_year           56
budget_adj            2614
revenue_adj           4840
dtype: int64
```

### 1.1.4 Data Cleaning

After determining the columns useful to our analysis, those that are not integral to our research questions are dropped. Afterwards, missing values and duplicates are dropped from the rows.

### 1.1.5 Drop Extraneous Columns

```
In [9]: # Dropping columns not to be used for analysis
df.drop(['cast', 'homepage', 'imdb_id', 'keywords', 'overview', 'tagline', 'vote_count'],
```

```
In [10]: # Checking to see if columns have been dropped
df.head()
```

```
Out[10]:
```

	id	popularity	director	runtime	\
0	135397	32.985763	Colin Trevorrow	124	
1	76341	28.419936	George Miller	120	
2	262500	13.112507	Robert Schwentke	119	
3	140607	11.173104	J.J. Abrams	136	

```

4  168259      9.335014      James Wan      137

                                genres \
0  Action|Adventure|Science Fiction|Thriller
1  Action|Adventure|Science Fiction|Thriller
2      Adventure|Science Fiction|Thriller
3  Action|Adventure|Science Fiction|Fantasy
4      Action|Crime|Thriller

                                production_companies  release_year \
0  Universal Studios|Amblin Entertainment|Legenda...      2015
1  Village Roadshow Pictures|Kennedy Miller Produ...      2015
2  Summit Entertainment|Mandeville Films|Red Wago...      2015
3      Lucasfilm|Truenorth Productions|Bad Robot      2015
4  Universal Pictures|Original Film|Media Rights ...      2015

      budget_adj  revenue_adj
0  1.379999e+08  1.392446e+09
1  1.379999e+08  3.481613e+08
2  1.012000e+08  2.716190e+08
3  1.839999e+08  1.902723e+09
4  1.747999e+08  1.385749e+09

```

Then we have to further remove columns that contain null values and that can't be replaced with mean because they're strings

```

In [11]: # Drop empty rows
         df.dropna(inplace = True)

In [12]: # Check to see if empty rows have been trimmed
         df.shape

Out[12]: (9807, 9)

In [13]: # Drop duplicate values in dataset
         df.drop_duplicates(inplace = True)

In [14]: df.head()

Out[14]:
      id  popularity  director  runtime \
0  135397  32.985763  Colin Trevorrow    124
1   76341  28.419936   George Miller    120
2  262500  13.112507  Robert Schwentke    119
3  140607  11.173104    J.J. Abrams    136
4  168259   9.335014    James Wan    137

                                genres \
0  Action|Adventure|Science Fiction|Thriller
1  Action|Adventure|Science Fiction|Thriller

```

```

2         Adventure|Science Fiction|Thriller
3   Action|Adventure|Science Fiction|Fantasy
4               Action|Crime|Thriller

           production_companies  release_year  \
0  Universal Studios|Amblin Entertainment|Legenda...  2015
1  Village Roadshow Pictures|Kennedy Miller Produ...  2015
2  Summit Entertainment|Mandeville Films|Red Wago...  2015
3           Lucasfilm|Truenorth Productions|Bad Robot  2015
4  Universal Pictures|Original Film|Media Rights ...  2015

      budget_adj  revenue_adj
0  1.379999e+08  1.392446e+09
1  1.379999e+08  3.481613e+08
2  1.012000e+08  2.716190e+08
3  1.839999e+08  1.902723e+09
4  1.747999e+08  1.385749e+09

```

Our Genres column contains multiple genres for each movie, we can then split the genres column, and convert the column from a string to a list containing the genres pertaining to each movie, in order for analysis to be easier.

```

In [15]: # columns to split by "/"
        split_columns = ['genres']

        # apply split function to each column of each dataframe copy
        for c in split_columns:
            df[c] = df[c].apply(lambda x: x.split("|"))

```

```

In [16]: df.head()

```

```

Out[16]:      id  popularity      director  runtime  \
0  135397   32.985763   Colin Trevorrow    124
1    76341   28.419936    George Miller    120
2  262500   13.112507  Robert Schwentke    119
3  140607   11.173104     J.J. Abrams    136
4  168259    9.335014     James Wan    137

           genres  \
0  [Action, Adventure, Science Fiction, Thriller]
1  [Action, Adventure, Science Fiction, Thriller]
2           [Adventure, Science Fiction, Thriller]
3  [Action, Adventure, Science Fiction, Fantasy]
4           [Action, Crime, Thriller]

           production_companies  release_year  \
0  Universal Studios|Amblin Entertainment|Legenda...  2015
1  Village Roadshow Pictures|Kennedy Miller Produ...  2015
2  Summit Entertainment|Mandeville Films|Red Wago...  2015

```



3	Lucasfilm Truenorth Productions Bad Robot	2015
4	Universal Pictures Original Film Media Rights ...	2015

	budget_adj	revenue_adj
0	1.379999e+08	1.392446e+09
1	1.379999e+08	3.481613e+08
2	1.012000e+08	2.716190e+08
3	1.839999e+08	1.902723e+09
4	1.747999e+08	1.385749e+09

Now that our data has been properly cleaned, and unused columns have been dropped, missing values and duplicate values have been removed, it's time to answer the research questions.

## Exploratory Data Analysis

### 1.1.6 Research Question 1 (Which genres are most popular from year to year?)

```
In [17]: # To begin exploring our data, we groupby release year first
popularity_mean = df.groupby(['release_year']).popularity.mean()
popularity_mean.head()
```

```
Out[17]: release_year
1960    0.458932
1961    0.430438
1962    0.465245
1963    0.502706
1964    0.421091
Name: popularity, dtype: float64
```


```
In [18]: # We then look at the summary statistics
popularity_mean.describe()
```

```
Out[18]: count    56.000000
mean         0.588379
std          0.147449
min          0.308457
25%          0.486578
50%          0.572578
75%          0.674149
max          1.135148
Name: popularity, dtype: float64
```

From our summary statistics, we can see that the max mean popularity is quite high compared to other values, this could be due to improperly inputted values. To continue our analysis, we'll look at a visualization of our data.

```
In [19]: # Popularity visualized
popularity_mean.hist()
```

Out[19]: <matplotlib.axes.\_subplots.AxesSubplot at 0x7fc076d40128>



We can see here that the average or mean popularity of all genres is skewed to the right.

```
In [20]: from collections import Counter

# Because the genres for every movie are basically sets
# We can use a "flattened list" approach for this question
flattened_genres = [elem for sublist in df['genres'] for elem in sublist]
fav_genre = Counter(flattened_genres).most_common()

# retrieving top 5 genres
fav_genre[:5]
```

```
Out[20]: [('Drama', 4369),
          ('Comedy', 3438),
          ('Thriller', 2747),
          ('Action', 2235),
          ('Romance', 1570)]
```

Now we can retrieve the top 5 genres over the years from 1960 -2015, and we can see that Drama comes out as the most popular movie over this entire period. This means that over this period of years, drama had the most consistently high popularity.

### 1.1.7 Research Question 2 (What kind of properties are associated with movies that have high revenues?)

In [21]: *#Return a query of all properties*

df.describe

Out[21]: <bound method NDFrame.describe of

	id	popularity	director	runtime
0	135397	32.985763	Colin Trevorrow	124
1	76341	28.419936	George Miller	120
2	262500	13.112507	Robert Schwentke	119
3	140607	11.173104	J.J. Abrams	136
4	168259	9.335014	James Wan	137
...	...	...	...	...
10861	21	0.080598	Bruce Brown	95
10862	20379	0.065543	John Frankenheimer	176
10863	39768	0.065141	Eldar Ryazanov	94
10864	21449	0.064317	Woody Allen	80
10865	22293	0.035919	Harold P. Warren	74

	genres \
0	[Action, Adventure, Science Fiction, Thriller]
1	[Action, Adventure, Science Fiction, Thriller]
2	[Adventure, Science Fiction, Thriller]
3	[Action, Adventure, Science Fiction, Fantasy]
4	[Action, Crime, Thriller]
...	...
10861	[Documentary]
10862	[Action, Adventure, Drama]
10863	[Mystery, Comedy]
10864	[Action, Comedy]
10865	[Horror]

	production_companies	release_year \
0	Universal Studios Amblin Entertainment Legenda...	2015
1	Village Roadshow Pictures Kennedy Miller Produ...	2015
2	Summit Entertainment Mandeville Films Red Wago...	2015
3	Lucasfilm Truenorth Productions Bad Robot	2015
4	Universal Pictures Original Film Media Rights ...	2015
...	...	...
10861	Bruce Brown Films	1966
10862	Cherokee Productions Joel Productions Douglas ...	1966
10863	Mosfilm	1966
10864	Benedict Pictures Corp.	1966
10865	Norm-Iris	1966

	budget_adj	revenue_adj
0	1.379999e+08	1.392446e+09
1	1.379999e+08	3.481613e+08
2	1.012000e+08	2.716190e+08

```

3      1.839999e+08  1.902723e+09
4      1.747999e+08  1.385749e+09
...      ...      ...
10861  0.000000e+00  0.000000e+00
10862  0.000000e+00  0.000000e+00
10863  0.000000e+00  0.000000e+00
10864  0.000000e+00  0.000000e+00
10865  1.276423e+05  0.000000e+00

```

```
[9806 rows x 9 columns]>
```

```

In [23]: # We return a query showing only top 50% revenues
top_movies = df.query('revenue_adj > revenue_adj.mean()')

# Determine how many unique entries we have in data
top_movies.info()

```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 2166 entries, 0 to 10848
Data columns (total 9 columns):
id                2166 non-null int64
popularity        2166 non-null float64
director          2166 non-null object
runtime           2166 non-null int64
genres            2166 non-null object
production_companies 2166 non-null object
release_year      2166 non-null int64
budget_adj        2166 non-null float64
revenue_adj       2166 non-null float64
dtypes: float64(3), int64(3), object(3)
memory usage: 169.2+ KB

```

Here we can just confirm how many entries out of the entire data frame fall into the higher revenue range.

```
In [24]: top_movies.describe()
```

```

Out[24]:

```

	id	popularity	runtime	release_year	budget_adj \
count	2166.000000	2166.000000	2166.000000	2166.000000	2.166000e+03
mean	31583.251616	1.592239	113.243767	1999.242382	5.969112e+07
std	59323.470090	1.795896	24.243409	12.464485	5.088598e+07
min	11.000000	0.010335	44.000000	1960.000000	0.000000e+00
25%	2112.500000	0.683560	98.000000	1992.000000	2.308905e+07
50%	9881.000000	1.120763	110.000000	2002.000000	4.607896e+07
75%	21817.750000	1.862250	124.000000	2009.000000	8.374198e+07
max	417859.000000	32.985763	705.000000	2015.000000	3.683713e+08

```

revenue_adj

```

```

count    2.166000e+03
mean     2.358479e+08
std      2.486913e+08
min      5.691199e+07
25%      8.974871e+07
50%      1.487309e+08
75%      2.779126e+08
max      2.827124e+09

```

We can see from this table and comparing it to the table description of all properties the variables which have a lower mean value against revenue, and those which have a higher mean value.

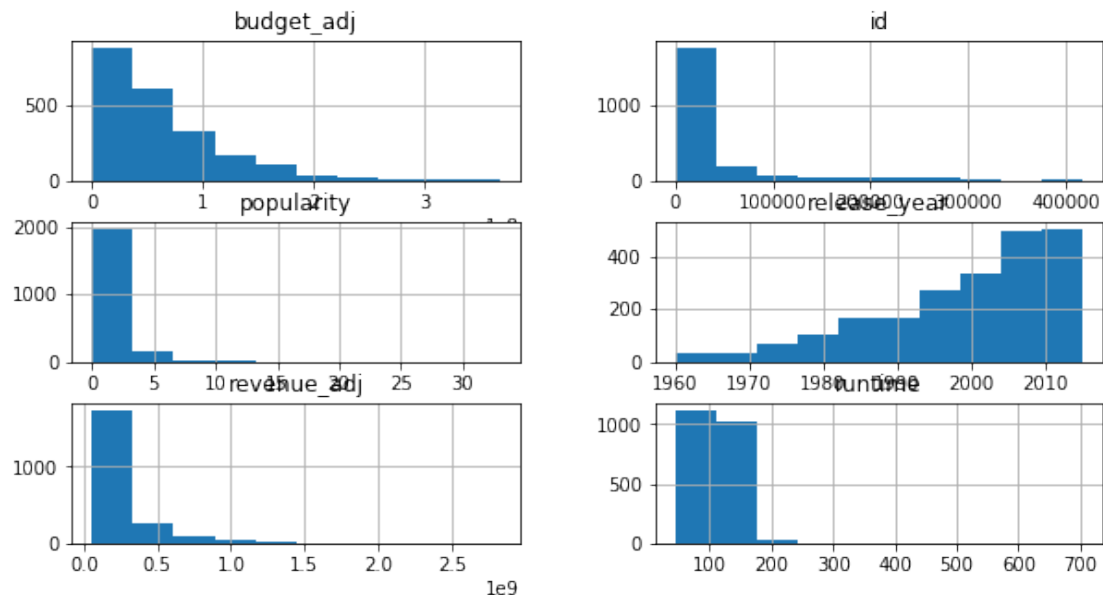
```
In [25]: # histogram to show variables
```

```
top_movies.hist(figsize= (10, 5))
```

```

Out[25]: array([[<matplotlib.axes._subplots.AxesSubplot object at 0x7fc0766e3278>,
<matplotlib.axes._subplots.AxesSubplot object at 0x7fc07668ee48>],
[<matplotlib.axes._subplots.AxesSubplot object at 0x7fc07663fe48>,
<matplotlib.axes._subplots.AxesSubplot object at 0x7fc076653588>],
[<matplotlib.axes._subplots.AxesSubplot object at 0x7fc07662ddd8>,
<matplotlib.axes._subplots.AxesSubplot object at 0x7fc07662de10>]],
dtype=object)

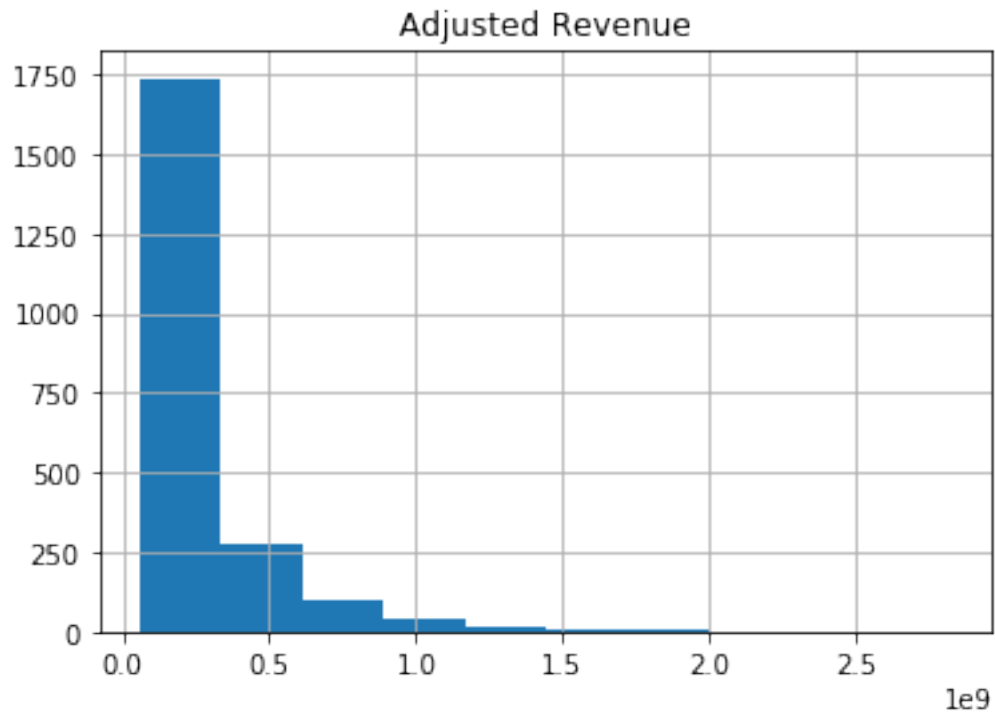
```



We can now take a look at our top earning movies. Then we create scatter plots to show the relationship between revenue and other variables to see if there is positive or negative correlation between the variables.

```
In [26]: top_movies.revenue_adj.hist()  
         plt.title("Adjusted Revenue")
```


```
Out[26]: Text(0.5,1,'Adjusted Revenue')
```



Now, we create scatter plots to visualize the relationship that exists between our dependent variable, the Revenue on the y axis, and the independent variables on the x axis.

```
In [27]: # Revenue against Budget  
         top_movies.plot(x = 'budget_adj', y = 'revenue_adj', kind = 'scatter')  
         plt.title("Adjusted Revenue against Budget")  
         plt.xlabel('Budget')  
         plt.ylabel('Revenue')
```


```
Out[27]: Text(0,0.5,'Revenue')
```



We can see that our adjusted revenue against budget seems to have minimal to no correlation.

```
In [28]: # Revenue against popularity
top_movies.plot(x = 'popularity', y = 'revenue_adj', kind = 'scatter')
plt.title("Adjusted Revenue against Popularity")
plt.xlabel('Popularity')
plt.ylabel('Revenue')
```

```
Out[28]: Text(0,0.5,'Revenue')
```




Likewise, we can see that Adjusted revenue has little to no correlation to popularity.

```
In [29]: # Revenue against runtime
top_movies.plot(x = 'runtime', y = 'revenue_adj', kind = 'scatter')
plt.title("Adjusted Revenue against Runtime")
plt.xlabel('Runtime')
plt.ylabel('Revenue')
```

```
Out[29]: Text(0,0.5,'Revenue')
```





After comparing our adjusted revenue against the three major quantitative independent variables, we find little to no correlation among them. This leads us to the conclusion that with the data we have, we're inconclusive on determining which variables have an impact, or are associated with high revenue movies.

## ## Conclusions

Our dataset was collected from Kaggle.com and is a compilation of movies of different genres from 1960 - 2015. The dataset contains a sizable number of missing values. To clean our data, we trimmed the dataset to only contain the necessary columns for our exploratory analysis, removed all duplicates found in our dataset and likewise dropped all null values from our dataset. In summary, in response to our first research question, we found that

In response to our second research question, we could not find any correlation, be it positive or negative between high revenue and any of the independent variables in the dataset. This might be due to the presence of outliers within said variables. Our statistical analysis showed that the high revenue generated by well-performing movies could not be traced back to any of our quantitative variables such as the budget, the popularity etc.

## 1.2 Limitations

A limitation in exploration results or findings in our dataset is that, the revenue variable for all movies might not be in US dollars. This is because data was collected from user input, and can't be fully verified.

Reference: Some part of our data analysis process was sourced and derived from [Kaggle.com](#)

```
In [30]: from subprocess import call  
         call(['python', '-m', 'nbconvert', 'Investigate_a_Dataset.ipynb'])
```

```
Out[30]: 0
```